

Introduction to Statistical Learning with Applications

CM3: Cross-validation, model selection,
and bias-variance

Pedro L. C. Rodrigues

Questions from the final exam from 2018

Consider the following Python script:

```
1  import numpy as np
2  import pandas as pd
3  import statsmodels.api as sm
4  np.random.seed(0)
5  # number of variables
6  pt = 201
7  # number of predictors
8  p = pt - 1
9  # sample size
10 n = 30 * p
11 # generate data
12 D = np.random.randn(n, pt)
13 df = pd.DataFrame(data=D)
14 df = df.rename(columns={0: 'Y'})
15 # do multiple linear regression
16 df['intercept'] = 1
17 model = sm.OLS(df['Y'], df.drop(columns='Y'))
18 results = model.fit()
19 print(results.summary())
```


- What does the script do?
- What is the true distribution of the random variable **Y** given the first 200 columns of **D** (X_1, X_2, \dots, X_{200})?
- Write an equation defining the model estimated by `model.fit()`. What is the difference between this model and the one defined above?
- Print `results.params` what's going on?

IN OUR PREVIOUS EPISODES...

Our current workflow


- 1 We want to estimate the values of Y based on predictors X_1, \dots, X_p
- 2 We are given a set of N examples of y_i with corresponding x_{i1}, \dots, x_{ip}
- 3 We estimate parameters $\hat{\beta}$ that minimize $\frac{1}{N} \sum_{i=1}^N (y_i - \beta^T x_i)^2$
- 4 We do statistical inference on the values of the $\hat{\beta}_j$
- 5 How can we **quantify the quality** of a model?

$$\downarrow x_i = \begin{bmatrix} x_{i1} \\ \vdots \\ x_{ip} \end{bmatrix}$$



- What makes a model **good**?
- Estimating the quality of a model
- Comparing and selecting models

What makes a model **good**?



Suppose our data is generated as


$$Y = \beta_0 + \sum_{i=1}^d \beta_i X^i + \varepsilon$$

we can run multiple linear regression with

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^d \end{bmatrix}$$

What's the best choice for d ?

What makes a model **good**?



Suppose our data is generated as


$$Y = \beta_0 + \sum_{i=1}^d \beta_i X^i + \varepsilon$$

we can run multiple linear regression with

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^d \end{bmatrix}$$

What's the best choice for d ?

What makes a model **good**?



Suppose our data is generated as

$$Y = \beta_0 + \sum_{i=1}^d \beta_i X^i + \varepsilon$$

we can run multiple linear regression with

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^d \end{bmatrix}$$

What's the best choice for d ?

What makes a model **good**? – The generalization error

Remember that our goal in regression was to minimize the generalization error:



$$\mathcal{L}(r) = \mathbb{E}_{Y,X} \left[\left(Y - r(X) \right)^2 \right]$$

Intuitively, we can expect this error to be decomposed into a few informative parts:

- The **irreducible error**: can we ever predict Y from X with zero prediction error? Probably not, since we always assume having some observation error ε in the data model.
- The **estimation bias**: when estimating the conditional expectation, we always have to choose a family of approximators, which may not always be sufficiently flexible.
- The **estimation variance**: for families of approximators that are too flexible, it might happen that for every small change in the observed data points, the estimated model changes.

What makes a model **good**?

Estimation variance. The same model fit on two slightly different datasets.



What makes a model **good**?



What makes a model **good**? – The bias-variance tradeoff


Remember that the data model is assumed to be $Y = r(X) + \varepsilon$ with $\text{Var}(\varepsilon) = \sigma^2$

We use a training dataset $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$ to estimate $\hat{r}_{\mathcal{D}}$

For a given new observation x of interest, we can write the bias-variance decomposition

$$\begin{aligned}\mathcal{L}(\hat{r}_{\mathcal{D}}(x)) &= \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \mid X = x \right] \\&= \mathbb{E}_{Y,X} \left[\left(Y - r(X) + r(X) - \hat{r}_{\mathcal{D}}(X) \right)^2 \mid X = x \right] \\&= \mathbb{E}_{Y,X} \left[\left(Y - r(X) \right)^2 \mid X = x \right] + \mathbb{E}_{Y,X} \left[\left(r(X) - \hat{r}_{\mathcal{D}}(X) \right)^2 \mid X = x \right] \\&= \underbrace{\mathbb{E}_{Y,X} \left[\left(Y - r(X) \right)^2 \mid X = x \right]}_{\text{Irreducible error } \sigma^2} + \underbrace{\left(\mathbb{E}_{Y,X} \left[r(X) - \hat{r}_{\mathcal{D}}(X) \mid X = x \right] \right)^2}_{\text{Squared estimation bias}} + \underbrace{\text{Var}(\hat{r}_{\mathcal{D}}(X) \mid X = x)}_{\text{Estimation variance}}\end{aligned}$$

- What makes a model **good**?



- Estimating the quality of a model

- Comparing and selecting models

Estimating the quality of a model

How can we estimate the generalization error of an estimated model in practice?

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right]$$



We don't
know the $p(Y, X)$

Estimating the quality of a model

How can we estimate the generalization error of an estimated model in practice?

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right] \approx \frac{1}{M} \sum_{i=1}^M \left(y_i - \hat{r}_{\mathcal{D}}(x_i) \right)^2 = L(\hat{r}_{\mathcal{D}}, \mathcal{X}) \quad \text{with } \mathcal{X} = \left\{ (y_i, x_i) \right\}_{i=1}^M$$



We don't
know the $p(Y, X)$

So we approximate it
with M data points

Note that for $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^N$ we have $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$

We want to minimize one
quantity but can only estimate
a proxy

Estimating the quality of a model

How can we estimate the generalization error of an estimated model in practice?

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right] \approx \frac{1}{M} \sum_{i=1}^M \left(y_i - \hat{r}_{\mathcal{D}}(x_i) \right)^2 = L(\hat{r}_{\mathcal{D}}, \mathcal{X}) \quad \text{with } \mathcal{X} = \left\{ (y_i, x_i) \right\}_{i=1}^M$$



We don't know the $p(Y, X)$

So we approximate it with M data points

Note that for $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^N$ we have $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$

We want to minimize one quantity but can only estimate a proxy

Question:

Can we say that $\mathcal{L}(\hat{r}_{\mathcal{D}}) \approx L(\hat{r}_{\mathcal{D}}, \mathcal{D})$?

ABSOLUTELY NOT!



In fact $\mathcal{L}(\hat{r}_{\mathcal{D}}) \geq L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ i.e. the true generalization error is larger than the estimated one

The optimism of the training error

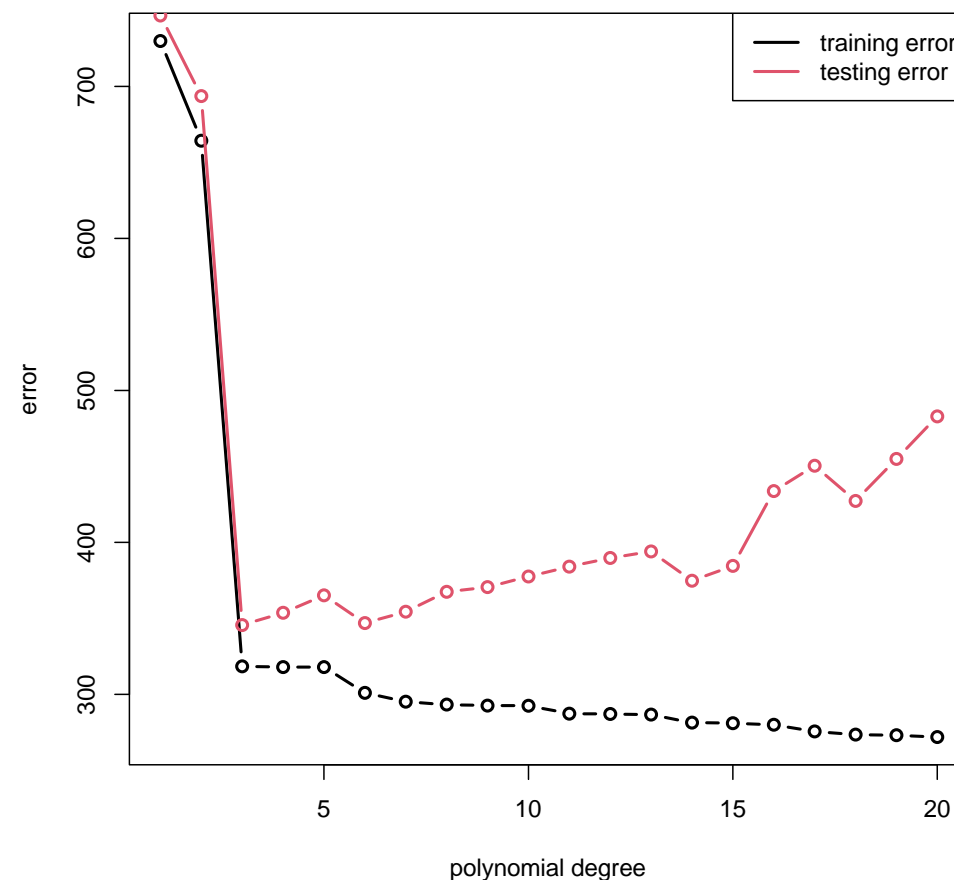
Different ways of seeing the problem:

- Intuitively, since $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$ then $L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ will always look small...
- Consider our example with polynomials
 - Increasing the degree improves training error
 - But the testing error grows!

$L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ is the **training** error

$L(\hat{r}_{\mathcal{D}}, \mathcal{X})$ with $\mathcal{X} \neq \mathcal{D}$ is the **testing** error

Error on unseen data



The optimism of the training error


Different ways of seeing the problem:

- Intuitively, since $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$ then $L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ will always look small...
- Consider our example with polynomials
 - Increasing the degree improves training error
 - But the testing error grows!
- Mathematical illustration on linear regression

$$\mathcal{D} = \left\{ (x_i, y_i) \right\}_{i=1}^{i=N} \quad \mathcal{D}' = \left\{ (x_i, y'_i) \right\}_{i=1}^{i=N} \quad Y = \beta_0 + \sum_{k=1}^p \beta_k X_k + \varepsilon$$

(same predictors but different observations)

$$\mathbb{E} \left[L(\hat{r}_{\mathcal{D}}, \mathcal{D}') \right] = \mathbb{E} \left[L(\hat{r}_{\mathcal{D}}, \mathcal{D}) \right] + \frac{2}{N} \sigma^2 (p + 1)$$




The optimism of the training error

Conclusion:

- We can only know if a model is good or not if we **correctly evaluate** its performance
- The objective function that we minimize when training a model (**training error**) is never the same as the one we are actually interested in minimizing (**test error**)

$$\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} \frac{1}{M} \sum_{i=1}^M \left(y_i - r(x_i) \right)^2$$

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{(X,Y)} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right]$$



Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

- **Strategy 1 : Single hold-out test point**

We fit a model on the first N-1 training samples, calling it $\hat{r}^{(-N)}$

Treat the last sample (X_N, Y_N) as a test sample and estimate generalization error as

$$\mathcal{L}(\hat{r}) \approx \left(y_N - \hat{r}^{(-N)}(x_N) \right)^2$$

The estimator is easy to calculate but we can expect a rather large variance.

Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

- **Strategy 2 : Leave one out cross-validation (LOOCV)**

We fit N models on N-1 training samples, holding out x_i at each time to get $\hat{r}^{(-i)}$

The test error is approximate as

$$\mathcal{L}(\hat{r}) \approx \frac{1}{N} \sum_{i=1}^N \left(y_i - \hat{r}^{(-i)}(x_i) \right)^2$$

The variance decreases, but the computational burden is much higher

Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

○ Strategy 3 : K-Fold cross-validation

Split the training dataset randomly into K folds so to have $\mathcal{D}_1 \cup \dots \cup \mathcal{D}_K = \mathcal{D}$

For $k=1, \dots, K$ fit a model $\hat{r}^{(-k)}$ on a training set except that excludes \mathcal{D}_k

$$\mathcal{L}(\hat{r}) \approx \frac{1}{K} \sum_{k=1}^K \left(\frac{1}{N_k} \sum_{i \in \mathcal{D}_k} \left(y_i - \hat{r}^{(-k)}(x_i) \right)^2 \right)$$

Take average
over folds




Estimate test error for each fold

Cross-validation

Strategy 3 : K-Fold cross-validation

Note that the errors calculated in each fold are not IID random variables, since the models share some of their data points.

↳ **Question:** What does this imply?



Using 

```
from sklearn.model_selection import KFold
from sklearn.datasets import fetch_california_housing
from sklearn.linear_model import LinearRegression
kf = KFold(n_splits=5)
dataset = fetch_california_housing()
lm = LinearRegression()
X = dataset.data
y = dataset.target
scores = []
for idx_train, idx_test in kf.split(X):
    X_train, y_train = X[idx_train], y[idx_train]
    X_test, y_test = X[idx_test], y[idx_test]
    lm.fit(X_train, y_train)
    scores.append(lm.score(X_test, y_test))
print(np.mean(scores))
```

See the documentation for more details

Cross-validation

But there are **several** other strategies!

<u>GroupKFold</u>	K-fold iterator variant with non-overlapping groups.
<u>GroupShuffleSplit</u>	Shuffle-Group(s)-Out cross-validation iterator.
<u>KFold</u>	K-Fold cross-validator.

<u>ShuffleSplit</u>	Random permutation cross-validator.
<u>StratifiedGroupKFold</u>	Stratified K-Fold iterator variant with non-overlapping groups.
<u>StratifiedKFold</u>	Stratified K-Fold cross-validator.
<u>StratifiedShuffleSplit</u>	Stratified ShuffleSplit cross-validator.
<u>TimeSeriesSplit</u>	Time Series cross-validator.

and more...

Cross-validation

(1) Example with the **categorical variables** from last week

```
from sklearn.model_selection import KFold, ShuffleSplit
from sklearn.linear_model import LinearRegression

regressor = LinearRegression()
cv = KFold(n_splits=4)
results_cv = cross_validate(
    regressor, X, y, cv=cv, scoring='neg_mean_squared_error')
print(f'MSE: {-results_cv["test_score"].mean():.2f}')

regressor = LinearRegression()
cv = ShuffleSplit(n_splits=4)
results_cv = cross_validate(
    regressor, X, y, cv=cv, scoring='neg_mean_squared_error')
print(f'MSE: {-results_cv["test_score"].mean():.2f}')
```

17.50

4.60

In [85]: df


Out[85]:

	life	rpm	brand
0	18.73	610	A
1	14.52	950	A
2	17.43	720	A
3	14.54	840	A
4	13.44	980	A
5	24.39	530	A
6	13.34	680	A
7	22.71	540	A
8	12.68	890	A
9	19.32	730	A
10	30.16	670	B
11	27.09	770	B
12	25.40	880	B
13	26.05	1000	B
14	33.49	760	B
15	35.62	590	B
16	26.07	910	B
17	36.78	650	B
18	34.95	810	B
19	43.67	500	B

What is going on? 😄

Cross-validation

(2) Example with **financial** time series: predict CVX's quotes based on other quotes



Cross-validation

(2) Example with **financial** time series: predict CVX's quotes based on other quotes

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import ShuffleSplit, cross_validate
X = quotes.drop(columns=["CVX"])
y = quotes["CVX"]
regressor = GradientBoostingRegressor()
cv = ShuffleSplit(n_splits=10)
results_cv = cross_validate(regressor, X, y, cv=cv)
print(f'Mean R2: {results_cv["test_score"].mean():.2f}')
```

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y}_i)^2}$$

0.95

This looks like an almost perfect prediction! 🤖 Does it sound right to you?

Cross-validation

(2) Example with **financial** time series: predict CVX's quotes based on other quotes

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import ShuffleSplit, cross_validate
X = quotes.drop(columns=["CVX"])
y = quotes["CVX"]
regressor = GradientBoostingRegressor()
cv = ShuffleSplit(n_splits=10)
results_cv = cross_validate(regressor, X, y, cv=cv)
print(f'Mean R2: {results_cv["test_score"].mean():.2f}')
```

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

0.95

This looks like an almost perfect prediction! 🤖 Does it sound right to you?

```
from sklearn.model_selection import TimeSeriesSplit
cv = TimeSeriesSplit(n_splits=10)
results_cv = cross_validate(regressor, X, y, cv=cv)
print(f'Mean R2: {results_cv["test_score"].mean():.2f}')
```

-3.10

Disappointing, but closer to reality...

- What makes a model **good**?

- Estimating the quality of a model

-  ○ Comparing and selecting models

Comparing and selecting models

Suppose we are given a dataset with p predictors.

We want to estimate a linear model with only a subset of them.

There are mainly **three strategies** for doing this properly:

- **Shrinkage** fits a model with all p predictors but using a modified loss function that drives some parameters to zero automatically. (Razan will talk about this)
- We can use **dimensionality reduction** techniques to project the p predictors to a lower-dimensional subspace. (We will see how to do this in CM4 and TP2)
- In **subset selection** we identify a subset of the predictors that seems the most adequate and then fit a model with them.

Comparing and selecting models

When p is large, testing all 2^p possible models can be very time consuming...

So we prefer to proceed greedily with e.g. **forward stepwise selection**

(1) Let \mathcal{M}_0 denote a model with no predictors (i.e. just the intercept)

(2) For $k = 0, \dots, p - 1$

(a) Consider all $p - k$ models that augment the predictors in \mathcal{M}_k by one extra predictor

(b) Choose the **best** among these $p - k$ models and call it \mathcal{M}_{k+1}

(3) Select single **best** model among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using cross-validated prediction error

Example: Consider the mtcars dataset – we want to predict **mpg**

Description:

The data was extracted from the 1974 _Motor Trend_ US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973-74 models).

Format:

A data frame with 32 observations on 11 variables.


[, 1]	mpg	Miles/(US) gallon
[, 2]	cyl	Number of cylinders
[, 3]	disp	Displacement (cu.in.)
[, 4]	hp	Gross horsepower
[, 5]	drat	Rear axle ratio
[, 6]	wt	Weight (lb/1000)
[, 7]	qsec	1/4 mile time
[, 8]	vs	V/S
[, 9]	am	Transmission (0 = automatic, 1 = manual)
[,10]	gear	Number of forward gears
[,11]	carb	Number of carburetors

OLS Regression Results


Dep. Variable:	mpg	R-squared:	0.869
Model:	OLS	Adj. R-squared:	0.807
Method:	Least Squares	F-statistic:	13.93
Date:	Fri, 27 Dec 2024	Prob (F-statistic):	3.79e-07
Time:	15:09:11	Log-Likelihood:	-69.855
No. Observations:	32	AIC:	161.7
Df Residuals:	21	BIC:	177.8
Df Model:	10		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
cyl	-0.1114	1.045	-0.107	0.916	-2.285	2.062
disp	0.0133	0.018	0.747	0.463	-0.024	0.050
hp	-0.0215	0.022	-0.987	0.335	-0.067	0.024
drat	0.7871	1.635	0.481	0.635	-2.614	4.188
wt	-3.7153	1.894	-1.961	0.063	-7.655	0.224
qsec	0.8210	0.731	1.123	0.274	-0.699	2.341
vs	0.3178	2.105	0.151	0.881	-4.059	4.694
am	2.5202	2.057	1.225	0.234	-1.757	6.797
gear	0.6554	1.493	0.439	0.665	-2.450	3.761
carb	-0.1994	0.829	-0.241	0.812	-1.923	1.524
intercept	12.3034	18.718	0.657	0.518	-26.623	51.229


Forward stepwise selection



Forward stepwise selection



Forward stepwise selection



– **Question 10: (credits to Berkeley CS-189)**

In the following statements, the word “bias” is referring to the bias-variance decomposition.

- (A) A model trained with N training points is likely to have lower variance than a model trained with $2N$ training points.
- (B) If my model is underfitting, it is more likely to have high bias than high variance.
- (C) Increasing the number of parameters (weights) in a model usually improves the test set accuracy.
- (D) None of the above.

– **Question 19: Bias-variance decomposition (credits to EPFL CS-433)**

Consider a regression model where data (x, y) is generated by input $x \in \mathbf{R}$ uniformly sampled between $[0, 1]$ and $y = x + \varepsilon$, where ε is random noise with mean 0 and variance 1. Two models are carried out for regression: model \mathcal{A} is a trained quadratic function $g_{\mathcal{A}}(x, \beta) = \beta_0 + \beta_1 x + \beta_2 x^2$ and model \mathcal{B} is a constant function $g_{\mathcal{B}}(x) = \frac{1}{2}$. Compared to model \mathcal{B} , model \mathcal{A} has

- (A) Higher bias, higher variance.
- (B) Lower bias, higher variance.
- (C) Higher bias, lower variance.
- (D) Lower bias, lower variance.

– **Question 22: Linear regression (credits to Berkeley CS-189)**

In linear regression, we model $p(y \mid \mathbf{x}) \sim \mathcal{N}(\beta^\top \mathbf{x} + \beta_0, \sigma^2)$. The irreducible error in this model is

- (A) σ^2
- (B) $\mathbb{E}[y \mid \mathbf{x}]$
- (C) $\mathbb{E}[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}]$
- (D) None of the above.